TECH DEBT

⚠ CRITICAL

MONOLITH DEPLOYMENT GRIDLOCK

Your monolithic codebase has grown so large that a full deployment now takes 4 hours and any failure requires a complete rollback of the entire system. The team can only release once a week. The feature backlog is months deep.

Deployment frequency is one of the four key DORA metrics — elite teams deploy on demand (multiple times per day); low performers deploy monthly or less. A deployment that takes 4 hours is not just an inconvenience; it is a structural ceiling on your entire engineering organisation's throughput. Strangler Fig is the standard pattern for incrementally decomposing a monolith: new functionality is built as separate services that sit in front of the monolith, gradually replacing it functionality by functionality without a risky big-bang rewrite. Meanwhile, trunk-based development, feature flags, and modular monolith architecture (well-separated internal modules even within a single deployable) can dramatically reduce deployment time and risk without full service extraction. The root cause is always accumulated coupling — every shortcut that bypassed modularity contributed to this moment.

Deployment frequency is a proxy for engineering health. If you cannot ship, you cannot compete.

TEAM DISCUSSION

How long does a deployment take end-to-end? What is the maximum safe release cadence your current architecture allows?