OPPORTUNITY

OPPORTUNITY

PERFORMANCE BENCHMARK GOES VIRAL

A respected independent blogger is running a head-to-head performance comparison in your space. Results go live next week. The last winner dominated the news cycle for a month and saw 3× trial signups. The last loser is still trying to recover.

Independent performance benchmarks are one of the highest-credibility forms of technical marketing — and one of the most unforgiving. Benchmarks typically measure throughput (requests per second), latency (p50, p95, p99 response times), concurrency handling, and resource efficiency. Teams with a continuous performance engineering practice — tracking these metrics in dashboards, setting performance budgets in CI, running regular load tests, and profiling hot paths proactively — can be confident in benchmark results because they already know the numbers. Teams who have not measured systematically often discover their p99 latency at the worst possible moment. The performance investment that pays off in a benchmark was made over the preceding sprints, not the week before results are published.

Benchmarks are unforgiving. You either performed consistently, or you did not. There is no narrative spin for a slow p99.

TEAM DISCUSSION

How confident are you in your performance profile right now? Have you measured and optimised your p95/p99 response times recently?