PERFORMANCE

ELEVATED

SLOW API UNDER LOAD

Response times have ballooned to 8 seconds during peak hours. Users are rage-tweeting, the app store rating dropped half a star overnight, and your largest customer just emailed asking what is going on.

Performance is a user-facing feature, not an infrastructure concern. Google research shows 53% of mobile users abandon a site that takes more than 3 seconds to load; Amazon found each 100ms of latency cost 1% of sales. The engineering response requires first understanding the bottleneck — distributed tracing (OpenTelemetry, Jaeger, Zipkin) shows exactly where latency accumulates across service boundaries. Common causes include N+1 database queries, missing indexes, synchronous calls that could be async, and missing caching. Performance budgets — defined SLOs for p95 and p99 response times — should be tracked in dashboards and trigger alerts when breached. Load testing in CI against a production-like environment catches regressions before they reach users.

Slow is the new down. Users tolerate outages more than chronic slowness. Performance is a feature, not an afterthought.

TEAM DISCUSSION

Do you have performance budgets and p95/p99 benchmarks defined? Do you know where the actual bottleneck is?