A/B TESTING / EXPERIMENTATION PLATFORMS

CONTINUOUS EXPERIMENTATION

Mature product teams don't guess — they experiment. Feature flags enable A/B testing, multivariate experiments, and gradual rollouts that convert product decisions from opinions into evidence. This requires fast deployments, granular targeting, reliable measurement, and the cultural permission to run experiments that might fail.

KEY BENEFITS

- ✓ Replaces opinion-driven decisions with statistically validated evidence
- ✓ Reduces risk of large feature launches — test with 5% before 100%
- ✓ Accelerates product learning — more experiments per quarter means faster improvement
- ✓ Creates organisational alignment — results are objective, not political

REAL-WORLD EXAMPLE

Microsoft runs over 10,000 concurrent experiments. Booking.com runs more than 1,000. This scale of experimentation is only possible when the deployment pipeline was built for it from the ground up.

PRACTICAL EXAMPLE

The team debates two checkout button designs. They ship both 50/50. After 48 hours and 4,000 sessions, version B shows a statistically significant 6.2% higher conversion rate. Decision made. No meeting required.

REFLECTION QUESTION

How many experiments is your team running right now? What infrastructure would you need to run 10× more?