SITE RELIABILITY ENGINEERING (SRE)

SLOS, SLIS AND ERROR BUDGETS

How much downtime is acceptable? Define it. Measure it. An error budget is the gap between 100% and your SLO — it's the team's licence to take risk. Burn the budget too fast and stability work takes priority over features, automatically, by policy, without debate.

KEY BENEFITS

- ✓ Eliminates the perpetual Dev vs Ops conflict over stability vs speed
- ✓ Provides an objective, data-driven basis for prioritisation decisions
- ✓ Aligns engineering investment with actual user experience, not gut feel
- ✓ Error budgets give teams permission to take risk when reliability is healthy

REAL-WORLD EXAMPLE

Google's SRE model introduced error budgets to eliminate the Dev/Ops conflict. When the budget is spent, everyone agrees on the priority — without politics. When healthy, teams have licence to ship fast.

PRACTICAL EXAMPLE

The team's SLO is 99.5% availability. By the 20th they've used 80% of their error budget. The next deployment is paused. They spend the remaining days fixing flaky tests and stabilising deploys. Next month the budget is healthy.

REFLECTION QUESTION

Do you have explicit uptime targets? Does your team know the current error budget balance?